

🎓 EDUCATION

Stony Brook University

Master of Science in Computer Science, Graduating Dec 2020, GPA: 3.85/4.0

Stony Brook, NY

Aug 2019 – Present

- Data Science, Natural Language Processing, Computer Vision, Big Data, Probability & Statistics.
- Teaching Assistant: Programming Abstractions (Fall '19), Software Engineering (Spring '20, Fall '21).

University of Mumbai

Bachelor of Engineering in Information Technology; GPA: 3.9/4.0

Mumbai, India

Aug 2011 – May 2015

- Data Structures & Algorithms, Discrete Maths, Software Engineering, Data Mining, Operating Systems

👛 EXPERIENCE

PlayStation, Software Engineer Intern

May 2020 – Aug 2020

- Built Deep Learning pipeline to improve perceived video quality by enhancing regions with text. Reduced inference time by ~75% for text detection under challenging constraints without reducing precision & recall.
- Evaluated multiple lightweight backbone networks, ResNet, PVANet, MobileNetV2, and techniques like quantization, post-processing optimization to further improve inference speed by 2x-4x.
- Developed a web dashboard using Flask, HTML & JavaScript, to help identify incorrect annotations which helped increase precision & recall of text detection by 1-2%.

JP Morgan Chase & Co., Associate Software Engineer

May 2017 – Jul 2019

- Redesigned server-side services to support Angular client & streaming of real-time data using WebSockets. Improved app performance with ~70% reduction in payload & boosted reliability using micro-services architecture with Spring Cloud.
- Developed non-intrusive ways to gather, visualize and analyze latency metrics using Elasticsearch, Logstash & Kibana, in the order of 100ms, across micro-services to identify bottlenecks in high throughput services.
- Optimized data collection mechanisms using REST APIs to track and compare client portfolio pre & post trade to generate risk reports on demand, increasing granularity in understanding firm's risk exposure.
- Designed a generic framework for database migration efforts to aid transition from PostgreSQL to MongoDB.
- Automated pre-release performance tests using CI/CD & build tools saving man-hours across teams.

LiveFiesta, Lead Android Developer

Jun 2016 – Jan 2017

- Led a team of 4 in the design & development of an Android application with an average rating of 4.5+ for customers to book tickets to events. Leveraged MVP architecture, Dependency Injection & TDD using RxJava, Dagger & Espresso.
- Reduced customer entry time to events by 50%, shortened queue lengths & cut losses due to fake ticket duplication & untracked re-entrants by developing a utility Android application to redeem tickets and track entrants.

TechGenium, Software Developer & Partner

Jun 2015 – May 2016

🎓 PUBLICATIONS

Querying Across Genres for Medical Claims in News

Empirical Methods in Natural Language Processing (EMNLP 2020)

Nov 2020

🔧 PROJECTS

- **Comment Toxicity Detection:** Achieved an AUC score of 0.98+ with creative pre-processing techniques coupled with Bi-GRU & BERT for multi-class classification of toxicity levels in Wikipedia comments.
- **Sentence Representations:** Understood sentence representations using Perturbation Analysis on sentiment analysis task with IMDb movie reviews using neural networks, DAN and GRU.
- **Summarizing Gameplay:** Leveraged Transformers & NLTK to summarize gameplay comments & sentiment analysis to inform e-sports players of crowd sentiment as feedback for immersive gameplay experience.

- **Video Action Classification:** Compared LSTM v/s SVM for action classification task on the UCF101 dataset. Leveraged Transfer Learning to compute features for 60000 video frames with limited compute resources.
- **3D Pose Estimation:** Estimated 3D pose co-ordinates of humans by regressing over their 2D co-ordinates only using Deep Neural Networks. Experimented with Kalman Filters for correcting noisy measurements due to faulty sensors.
- **Natural Scene Classification:** Classified scenes using Transfer Learning & Convolutional Neural Networks (AlexNet, VGG) with 89% accuracy. Used features from pre-trained networks to learn a Linear SVC for the same classification task.
- **Chess Rating Prediction:** Evaluated Random Forests, Gradient Boosting over novel features extracted from moves in 100k chess games to predict Elo ratings. Feature extraction run on distributed nodes using OpenMP for 15x speed.
- **Understanding Infant Mortality:** Applied Linear Regression to suggest priority actions to reduce Infant Mortality Rate based on 16.8GB of health & social records of 3M women using Dask & Apache Spark for parallel computation.
- **HoldingWilley** 🍏: iOS app for displaying real-time scores, stats & analysis of live cricket matches.
- **WaveView** 🔄 🌐: Android/Java library for rendering/animating sinusoidal waves with 10k+ downloads.

⚡ TECHNICAL SKILLS

- **Languages:** Proficient in Python & Java, Familiar with C & C++, JavaScript, SQL
- **Frameworks & Libraries:** Spring, Android, iOS, Flask, TensorFlow, PyTorch, scikit-learn, OpenCV, NLTK, Pandas, Numpy, Kafka, Docker, Kubernetes, Dask, Hadoop, Spark, MySQL, PostgreSQL, MongoDB
- **Build Tools:** Git, Gradle, Maven, Jenkins, Bash, Linux, Google Cloud Platform(GCP)